Minimum Message Length and Kolmogorov Complexity

C. S. Wallace and D. L. Dowe

Presentation by Alkis Gkotovos

Overview

Kolmogorov complexity

Quantify complexity of binary strings via Turing Machines (early '60s)

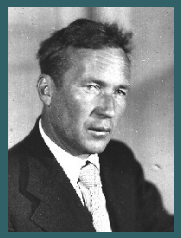
A. Kolmogorov

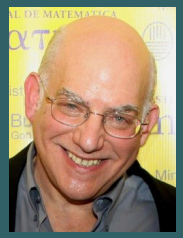
G. Chaitin

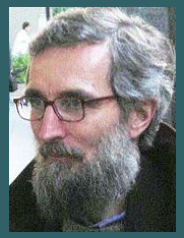
P. Martin-Löf

?

Universal induction

Define algorithmic probability via Turing Machines and use it for induction (early '60s)


R. Solomonoff

MML/MDL

Infer a hypothesis about the data via two-part coding (late '60s and '70s)


C. Wallace


J. Rissanen

Minimum Message/Description Length

Data string

Encode x using a *two-part* scheme

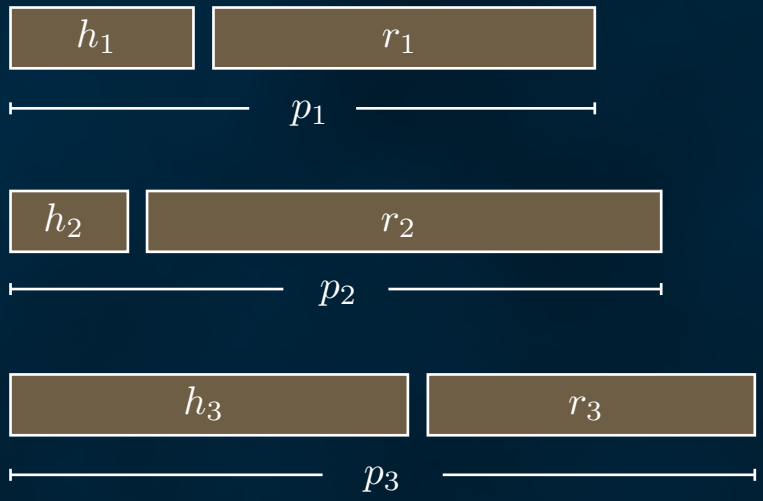

Pick the hypothesis that results in the minimum encoding length

$$l(p_i) = l(h_i) + l(r_i) = -\log_2(p_H(h_i)) - \log_2(p_X(x \mid h_i))$$

Turing Machines

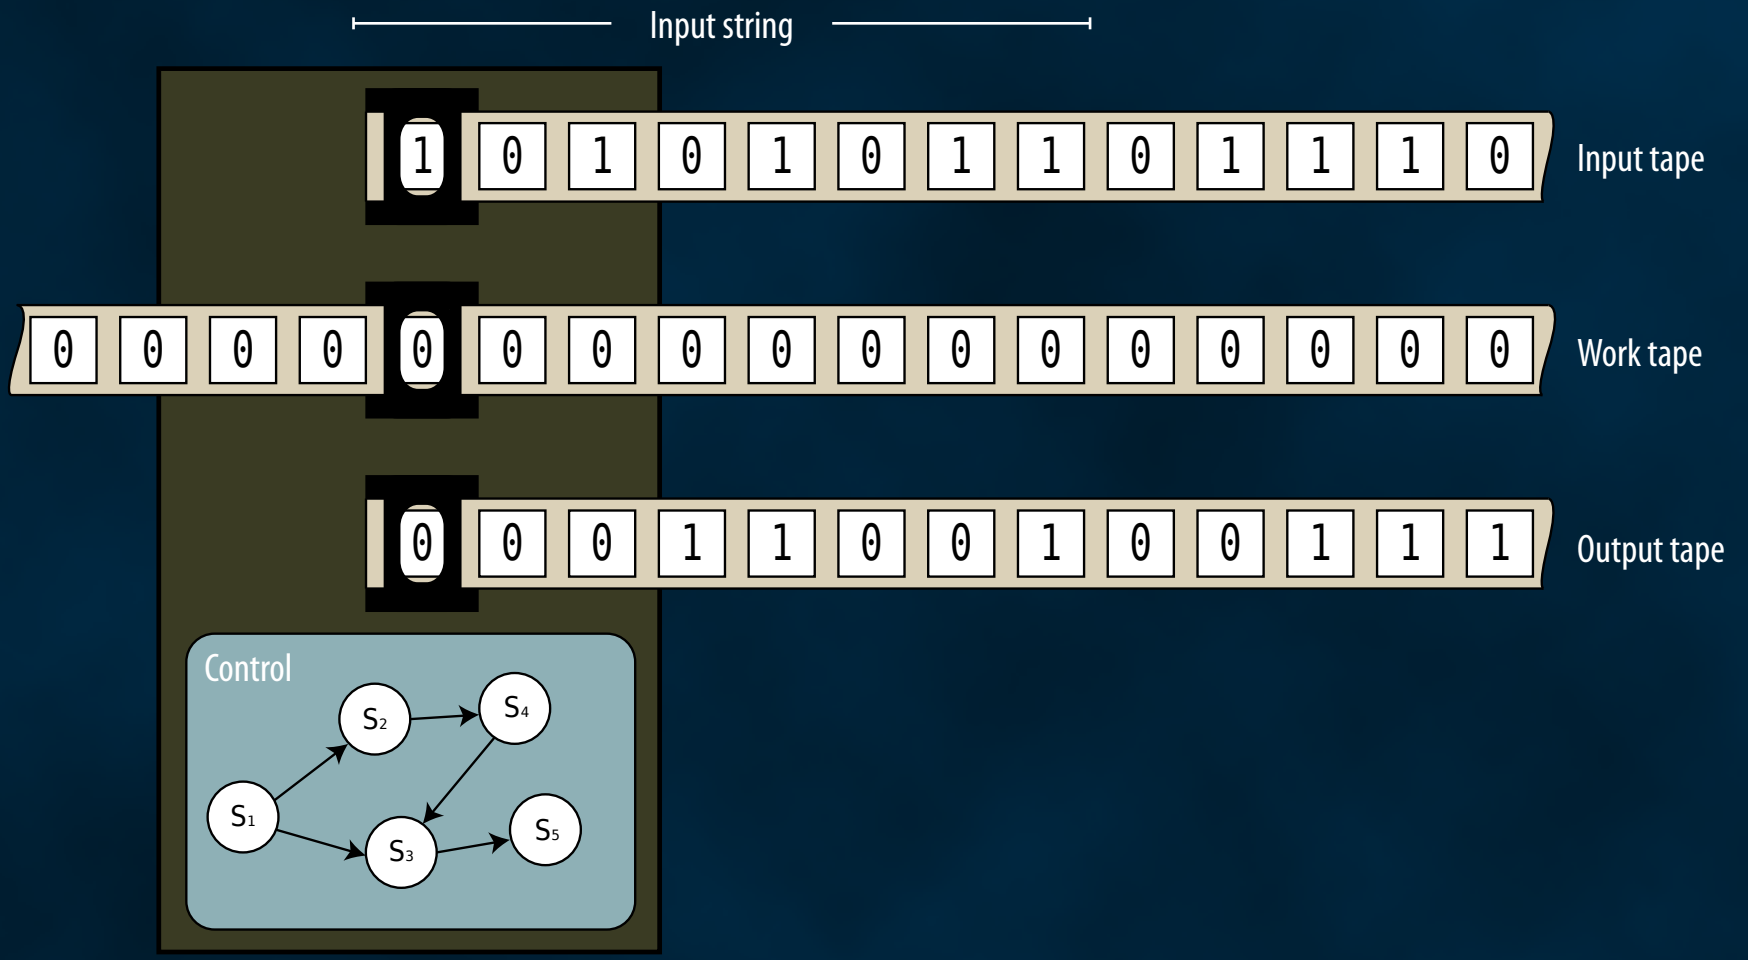

Data & Hypotheses

Data string x is a representation of observational data from a real world phenomenon

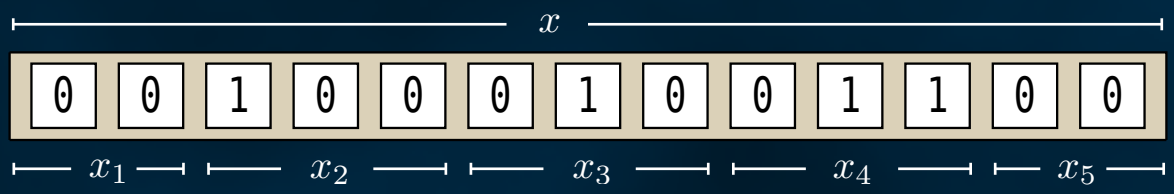

$$L = \{00, 100, 010, 011\}$$

- "Sentences" $x_i \in L$, where L is a prefix-free set (data "language")
- Distinct sentences represent distinct real-world facts
- Sentences are conditionally independent given full knowledge of the phenomenon
- Strings are invariant to sentence permutation

Data & Hypotheses

Hypothesis Q is a (computable) probability distribution over L

Conditional independence of sentences implies

$$x = x_1 \dots x_n \Rightarrow Q(x) = Q(x_1) \times \dots \times Q(x_n)$$

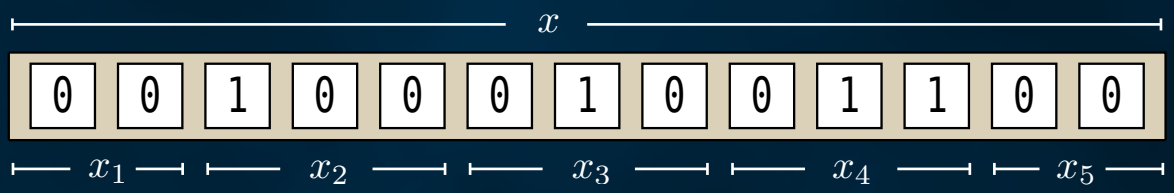

$$L = \{00, 100, 010, 011\}$$

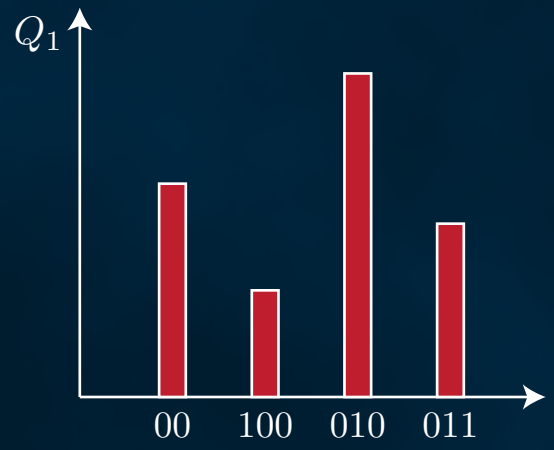

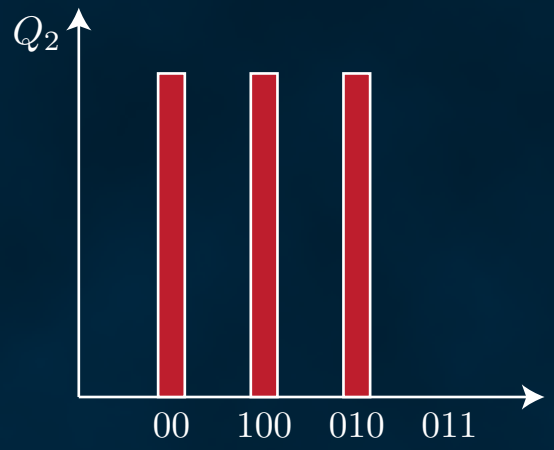

Two-part encoding

How do we acquire a hypothesis-based encoding of data in the Algorithmic Complexity framework?

Idea

- Use conditional Kolmogorov complexity

$$K_T(x \mid y) = \min\{l(p) \mid T(\langle y, p\rangle) = x\}$$

  and interpret y as hypothesis and x as data

- Corresponding conditional algorithmic probability

$$P_T(x \mid y) = 2^{-K_T(x|y)}$$

Problem

Probability can never be 0, i.e. Popper-falsification not possible, because

$$K(x \mid y) < K(x) + O(1) \Rightarrow P_K(x \mid y) > P_K(x) + O(1)$$

Why? Hypothesis y acts as "extra info", instead of assertively

Proposal

- Have hypothesis be a prefix of input string p
- Force intended two-part encoding by imposing conditions on p

Two-part encoding

Input p is an acceptable MML message encoding data string x, if

1)	$T(p) = x$	p encodes x
2)	$l(p) < l(x)$	some compression is achieved
3)	$p = qr$	two-part encoding
4)	$T(q) = \epsilon$	hypothesis q is does not determine data
5)	$T_q(rs) = xT_q(s)$	reading r does not alter the state of T
6)	$l(r) < K_T(x)$	hypothesis q is "significant"
7)	$x = x_1 \dots x_n \Rightarrow \begin{cases} r = r_1 \dots r_n \\ T_q(r_i) = x_i,\ i = 1 \dots n \end{cases}$	conditionally independent sentences
8)	$\begin{matrix} x' = x^{(1)}x^{(2)} \\ j' = j^{(1)}j^{(2)} \end{matrix} \Rightarrow \begin{matrix} T_q(j^{(1)}) = x^{(1)},\ j^{(1)} < K_T(x^{(1)}) \\ T_q(j^{(2)}) = x^{(2)},\ j^{(2)} < K_T(x^{(2)}) \end{matrix}$	hypothesis q is "general"
9)	No prefix of q satisfies all the above conditions	all of q is required

Two-part encoding

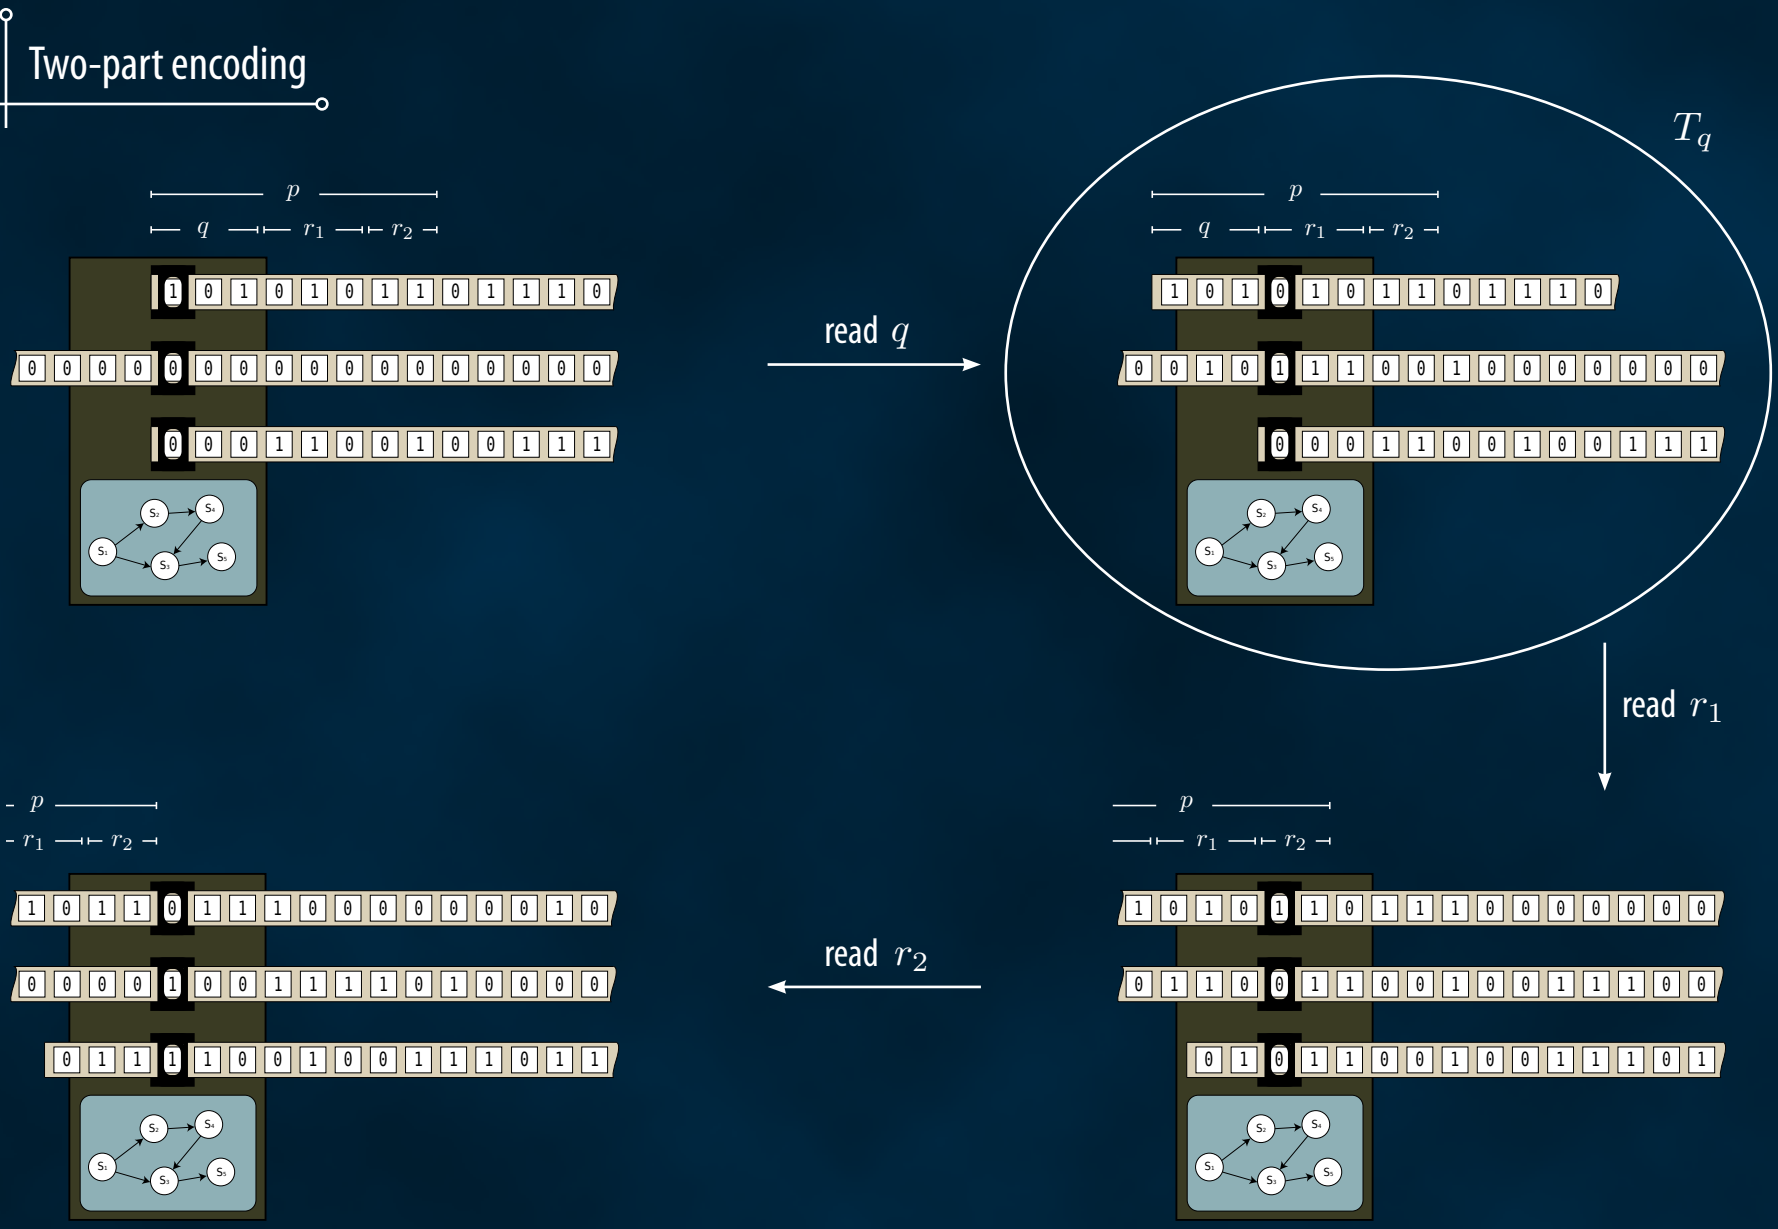

Two-part encoding

- The division of p into q and r is unique

- In what way exactly does hypothesis string q affect T?

Remember $T \xrightarrow{q} T_q$

T_q is a decoder of "second parts"

$$T_q : S \to W$$

Code words

$$S = \{r_i \in \{0,1\}^* \mid T_q(r_i) \in L\}$$

Subset of L that is coded

$$W = \{x_i \in L \mid \exists r_i \in S : T_q(r_i) = x_i\}$$

In fact, T_q decodes a prefix code (why?)

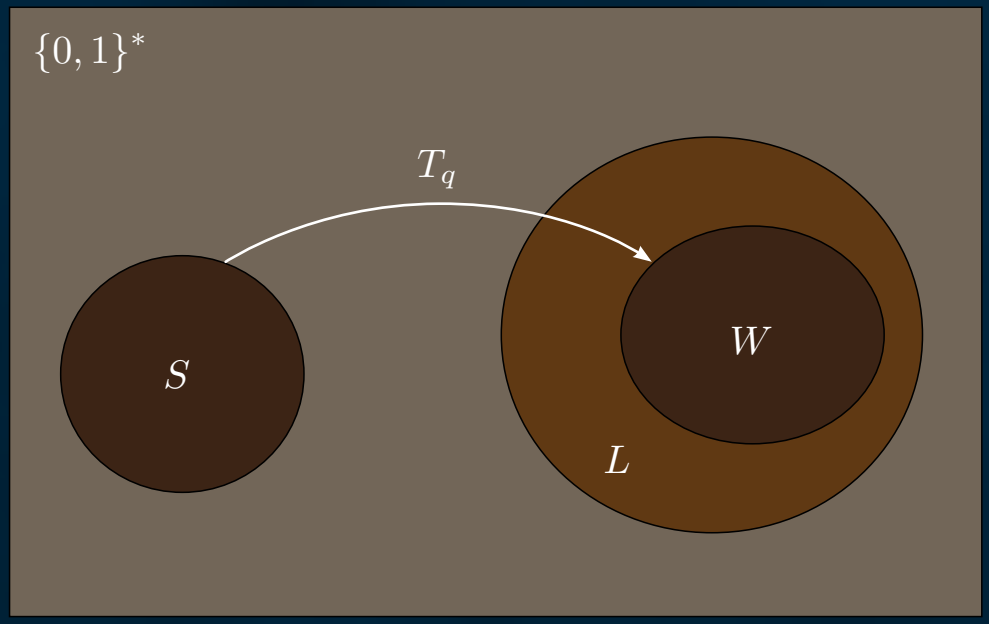

Two-part encoding

- What is the hypothesis (probability distribution) Q implied by hypothesis string q?

$$Q(x_i) = \begin{cases} 2^{-l(p)} & , \text{ if } p \text{ is a shortest codeword for sentence } x_i \in L \\ 0 & , \text{ if there is no codeword for sentence } x_i \in L \end{cases}$$

Because of prefix code

$$\sum_{x_i \in L} Q(x_i) = \sum_{x_i \in W} 2^{-l(p)} \overset{\text{Kraft}}{\leq} 1$$

- In this setting, hypotheses are falsifiable:

2) $l(p) < l(x) \Rightarrow l(r) < l(x)$

If Q assigns low probability (eq. high codeword length) to a sentence x_i, then adding enough such sentences to the data string will violate the above condition and falsify the hypothesis

Can Q assign lower codeword length to every sentence?
(L is a complete prefix code for "data facts")

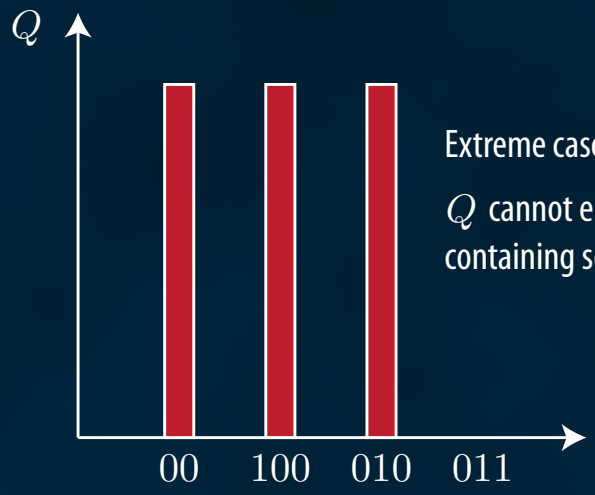

Extreme case:

Q cannot encode *any* data string containing sentence 011.

Two-part encoding

- What do we "pay" for enforcing a two-part encoding scheme?

Shortest acceptable MML input string: $M_T(x)$
Shortest unconstrained string: $K_T(x)$

with $M_T(x) \;\; ? \;\; K_T(x)$

$$
\begin{aligned}
M_T(x) - K_T(x) &= l(q) + l(r) - K_T(x) \\
&= K_T(Q) - \log_2(Q(x)) - K_T(x) \\
&= -\log_2\left(\frac{P_T(Q)Q(x)}{P_T(x)}\right) \\
&\approx -\log_2(\Pr(Q \mid x))
\end{aligned}
$$

$$P_T(x) = 2^{-K_T(x)}$$

Finding the shortest MML string is like MAP, where $P_T(Q)$ plays the role of the prior

The log posterior odds ratio of two hypotheses is

$$\log_2\left(\frac{\Pr(Q_1 \mid x)}{\Pr(Q_2 \mid x)}\right) = l(p_1) - l(p_2)$$

where p_1 and p_2 are shortest input strings for their respective hypotheses