



Minimum Message Length and Kolmogorov Complexity

C. S. Wallace and D. L. Dowe



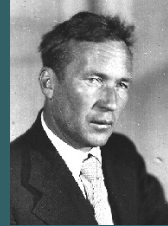
Induction



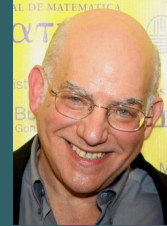
- Use some principle to guide inductive process, e.g. Occam's razor (bias towards "simplicity")
 - How do we quantify complexity?
 - Probability theory (e.g. AIC, BIC)
 - Information theory (e.g. MDL, MML)
 - Algorithmic complexity
- } intimately related (probability \longleftrightarrow codeword length)

Kolmogorov complexity

Quantify complexity of binary strings via Turing Machines (early '60s)



A. Kolmogorov



G. Chaitin



P. Martin-Löf



Universal induction

Define algorithmic probability via Turing Machines and use it for induction (early '60s)



R. Solomonoff

MML/MDL

Infer a hypothesis about the data via two-part coding (late '60s and '70s)





C. Wallace



J. Rissanen

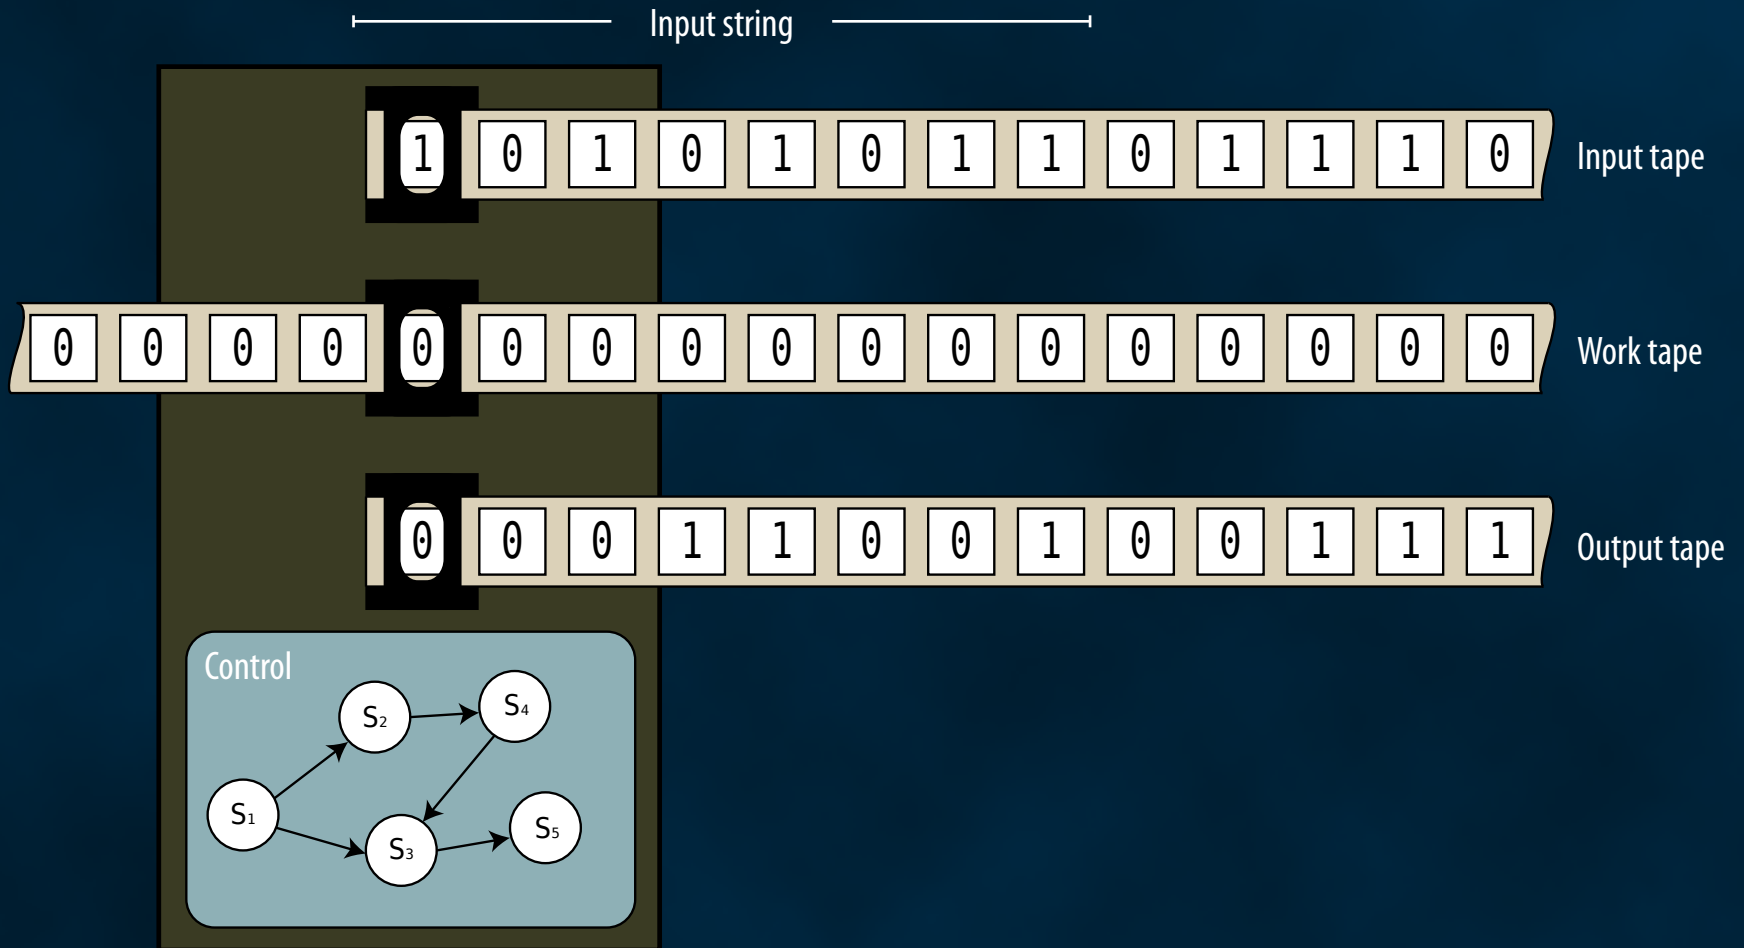
Minimum Message/Description Length

Data string

Encode x using a *two-part* scheme

Pick the hypothesis that results in the minimum encoding length


$$l(p_i) = l(h_i) + l(r_i) = -\log_2(p_H(h_i)) - \log_2(p_X(x | h_i))$$



Prefix TM

$$T(p) = x \text{ if}$$

- T reads all of the input p
- T writes x to the output
- T halts without reading/writing anything else



$$T(101010) = 011$$

T is a decoder of a prefix code (why prefix?)

Universal (prefix) TM = TM that can emulate any other (prefix) TM, e.g. $T(\langle i, p \rangle) = T_i(p)$

(Prefix) Kolmogorov complexity of x = The length of the shortest input string required to produce x

$$K_T(x) = \min\{l(p) \mid T(p) = x\}$$

Definition dependent on T ! Does this make any sense?

- Invariance theorem: Choose a *universal* prefix TM (among a special class) \longrightarrow as good as any TM (up to a constant)
- Intuition: You can do (almost) as good as any TM by writing a “compiler” for it
- Caveat 1: Kolmogorov complexity is not computable (approximate, e.g. via compressor)
- Caveat 2: Constants can be large (more on this later)


Can define a “probability” measure from Kolmogorov complexity

$$P_T(x) = 2^{-K_T(x)} \quad \text{with} \quad \sum_x P_T(x) \leq 1 \quad (\text{why?})$$

Why don't we normalize?

(Hereafter: “probability” \longleftrightarrow semimeasure)

Data string x is a representation of observational data from a real world phenomenon




$$L = \{00, 100, 010, 011\}$$

- “Sentences” $x_i \in L$, where L is a prefix-free set (data “language”)
- Distinct sentences represent distinct real-world facts
- Sentences are conditionally independent given full knowledge of the phenomenon
- Strings are invariant to sentence permutation



Hypothesis Q is a (computable) probability distribution over L

Conditional independence of sentences implies

$$x = x_1 \dots x_n \Rightarrow Q(x) = Q(x_1) \times \dots \times Q(x_n)$$



$$L = \{00, 100, 010, 011\}$$



How do we acquire a hypothesis-based encoding of data in the Algorithmic Complexity framework?

Idea

- Use conditional Kolmogorov complexity

$$K_T(x \mid y) = \min\{l(p) \mid T(\langle y, p \rangle) = x\}$$

and interpret y as hypothesis and x as data

- Corresponding conditional algorithmic probability

$$P_T(x \mid y) = 2^{-K_T(x \mid y)}$$

Problem

Probability can never be 0, i.e. Popper-falsification not possible, because

$$K_T(x \mid y) < K_T(x) + O(1) \Rightarrow P_T(x \mid y) > P_T(x) + O(1)$$

Why? Hypothesis y acts as “extra info”, instead of assertively

Proposal


- Have hypothesis be a prefix of input string p
- Force intended two-part encoding by imposing conditions on p

Input p is an acceptable MML message encoding data string x , if


- 1) $T(p) = x$ p encodes x
- 2) $l(p) < l(x)$ some compression is achieved
- 3) $p = qr$ two-part encoding
- 4) $T(q) = \epsilon$ hypothesis q is does not determine data
- 5) $T_q(rs) = xT_q(s)$ reading r does not alter the state of T
- 6) $l(r) < K_T(x)$ hypothesis q is "significant"
- 7) $x = x_1 \dots x_n \Rightarrow \begin{cases} r = r_1 \dots r_n \\ T_q(r_i) = x_i, \ i = 1 \dots n \end{cases}$ conditionally independent sentences
- 8) $\begin{matrix} x' = x^{(1)}x^{(2)} \\ j' = j^{(1)}j^{(2)} \end{matrix} \Rightarrow \begin{matrix} T_q(j^{(1)}) = x^{(1)}, \ j^{(1)} < K_T(x^{(1)}) \\ T_q(j^{(2)}) = x^{(2)}, \ j^{(2)} < K_T(x^{(2)}) \end{matrix}$ hypothesis q is "general"
- 9) No prefix of q satisfies all the above conditions all of q is required

TWO-PART ENCODING

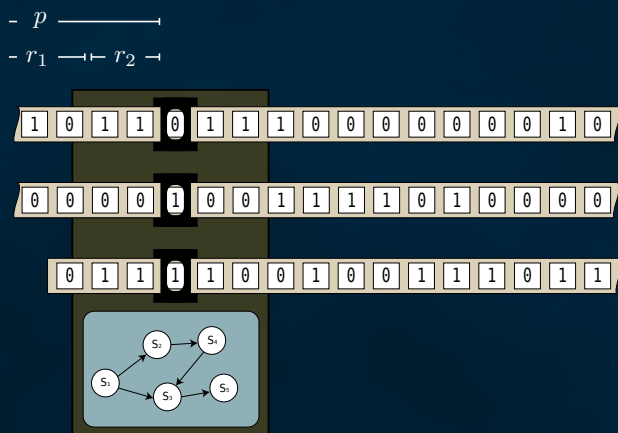
CONDITIONS



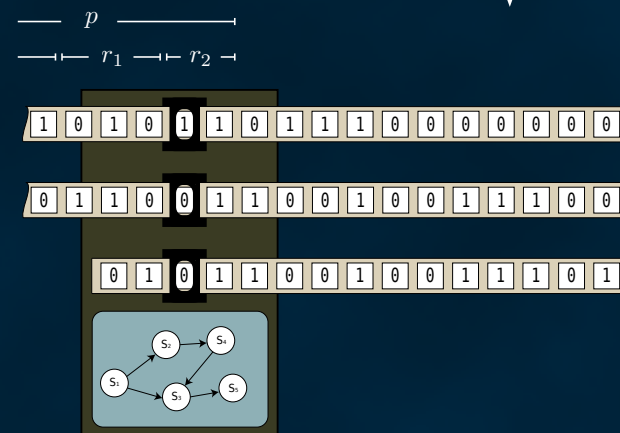
read q



read r_1



read r_2



- The division of p into q and r is unique
- In what way exactly does hypothesis string q affect T ?

Remember $T \xrightarrow{q} T_q$

T_q is a decoder of “second parts”

$$T_q : S \rightarrow W$$


Code words

$$S = \{r_i \in \{0, 1\}^* \mid T_q(r_i) \in L\}$$

Subset of L that is coded

$$W = \{x_i \in L \mid \exists r_i \in S : T_q(r_i) = x_i\}$$

In fact, T_q decodes a prefix code (why?)



- What is the hypothesis (probability distribution) Q implied by hypothesis string q ?

$$Q(x_i) = \begin{cases} 2^{-l(p)} & , \text{ if } p \text{ is a shortest codeword for sentence } x_i \in L \\ 0 & , \text{ if there is no codeword for sentence } x_i \in L \end{cases}$$

Because of prefix code


$$\sum_{x_i \in L} Q(x_i) = \sum_{x_i \in W} 2^{-l(p)} \stackrel{\text{Kraft}}{\leq} 1$$

- In this setting, hypotheses are falsifiable:

$$2) \quad l(p) < l(x) \Rightarrow l(r) < l(x)$$

If Q assigns low probability (eq. high codeword length) to a sentence x_i , then adding enough such sentences to the data string will violate the above condition and falsify the hypothesis

Can Q assign lower codeword length to every sentence?
(L is a complete prefix code for "data facts")



- What do we “pay” for enforcing a two-part encoding scheme?

Shortest acceptable MML input string: $M_T(x)$ with $M_T(x) \leq K_T(x)$
 Shortest unconstrained string: $K_T(x)$

$$\begin{aligned}
 M_T(x) - K_T(x) &= l(q) + l(r) - K_T(x) \\
 &= K_T(Q) - \log_2(Q(x)) - K_T(x) \\
 &= -\log_2 \left(\frac{P_T(Q)Q(x)}{P_T(x)} \right) \\
 &\approx -\log_2(\Pr(Q \mid x))
 \end{aligned}$$

$$P_T(x) = 2^{-K_T(x)}$$

Finding the shortest MML string is like MAP, where $P_T(Q)$ plays the role of the prior

The log posterior odds ratio of two hypotheses is

$$\log_2 \left(\frac{\Pr(Q_1 \mid x)}{\Pr(Q_2 \mid x)} \right) = l(p_2) - l(p_1)$$

where p_1 and p_2 are shortest input strings for their respective hypotheses

- Solomonoff: *Truly* Bayesian universal induction

Universal prior = Bayesian mixture over all possible “theories” (semicomputable semimeasures)

- Hutter: Many (all?) interesting problems \longrightarrow Sequence prediction \longrightarrow (Universal) induction + Decision theory

- Why would we ever need to pick one “theory”?!

- Universal induction is uncomputable
- Even constrained version is *highly* infeasible (right now, for real problems)
 - 1) Encode all human knowledge into a huge string to use as prior
 - 2) Compute Bayesian posterior for each event we want to predict
- General theories are compact, efficient (but imperfect) summaries of prior knowledge
- Humans prefer to understand the world as general theories, instead of mixtures

[...] Bayes/Solomonoff is the Gold standard for prediction, but MML/MDL is (often) a good “approximation/simplification” for explanation and understanding.


M. Hutter

- **Compromise:** Instead of a single theory, retain the few best of them

- Results in Kolmogorov complexity are almost always “up to a constant”
- Constants correspond to length of “compiler” and can be large
- Remember $\log_2 \left(\frac{\Pr(Q_1 | x)}{\Pr(Q_2 | x)} \right) = l(p_2) - l(p_1)$

Example $l(p_2) - l(p_1) = 10 \Rightarrow \Pr(Q_1 | x) = 1024 \times \Pr(Q_2 | x)$

- In practice, one has to choose priors very carefully in order to avoid unwanted biases
(authors claim that MML school is more considerate in this regard than Kolmogorov and MDL ones)
- Maybe time to change the paradigm?



THE END

T	H	A	N	K		Y	O	U	!
---	---	---	---	---	--	---	---	---	---